

Keegan Harris

Curriculum Vitae - June 2026

Email: keegan.harris@berkeley.edu

Web: keeganharris.github.io

Employment

University of California, Berkeley

2026 – Present

Postdoctoral Researcher

- Postdoc at the [Berkeley AI Research Lab](#) with [Nika Haghtalab](#) and [Michael Jordan](#)
- Honors: [Simons-Berkeley Research Fellowship](#), [ICML Spotlight paper](#)

Education

Carnegie Mellon University

- Doctor of Philosophy in Machine Learning 2020 – 2025
- Master of Science in Machine Learning 2022
- Thesis: [Incentives, Information, and Uncertainty in Data-Driven Decision Making](#)
- Committee: [Nina Balcan](#) (advisor), [Steven Wu](#) (advisor), [Michael Jordan](#), [Tuomas Sandholm](#), [Alex Slivkins](#)
- Honors: [NDSEG Fellowship](#), named a [2024 Rising Star in Data Science](#)

Pennsylvania State University

- Bachelor of Science in Computer Science, *summa cum laude* 2016 – 2019
- Bachelor of Science in Physics, *summa cum laude* 2016 – 2019
- Minor in Mathematics
- Honors: [Evan Pugh Scholar Award](#), [Phi Beta Kappa](#) (early induction), [Sigma Pi Sigma](#) (early induction)

Publications

1. **Keegan Harris** and Aleksandrs Slivkins. *Should You Use Your Large Language Model to Explore or Exploit?* Conference on Uncertainty in Artificial Intelligence (UAI), 2026.
2. Maria-Florina Balcan*, Kiriaki Fragkia*, and **Keegan Harris***. *Learning in Structured Stackelberg Games*. International Conference on Machine Learning (ICML), 2026. **Spotlight presentation (top 2.2%)**.
3. Maria-Florina Balcan*, Martino Bernasconi*, Matteo Castiglioni*, Andrea Celli*, **Keegan Harris***, and Zhiwei Steven Wu*. *Nearly-Optimal Bandit Learning in Stackelberg Games with Side Information*. International Conference on Learning Representations (ICLR), 2026.
4. Daniel Ngo*, **Keegan Harris***, Anish Agarwal, Vasilis Syrgkanis, and Zhiwei Steven Wu. *Incentive-Aware Synthetic Control: Accurate Counterfactual Estimation via Incentivized Exploration*. Transactions on Machine Learning Research (TMLR), 2026.
5. Akshay Krishnamurthy, **Keegan Harris**, Dylan J. Foster, Cyril Zhang, and Aleksandrs Slivkins. *Can Large Language Models Explore In-Context?* Neural Information Processing Systems (NeurIPS), 2024.
6. **Keegan Harris**, Zhiwei Steven Wu, and Maria-Florina Balcan. *Regret Minimization in Stackelberg Games with Side Information*. Neural Information Processing Systems (NeurIPS), 2024.
7. **Keegan Harris**, Anish Agarwal, Chara Podimata, and Zhiwei Steven Wu. *Strategyproof Decision-Making in Panel Data Settings and Beyond*. ACM SIGMETRICS International Conference on Measurement and Modeling of Computer Systems (SIGMETRICS), 2024.
8. Anish Agarwal*, **Keegan Harris***, Justin Whitehouse*, and Zhiwei Steven Wu*. *Adaptive Principal Component Regression with Applications to Panel Data*. Neural Information Processing Systems (NeurIPS), 2023.
9. Mikhail Khodak*, Ilya Osadchiy*, **Keegan Harris**, Maria-Florina Balcan, Kfir Levy, Ron Mier, and Zhiwei Steven Wu. *Meta-Learning Adversarial Bandit Algorithms*. Neural Information Processing Systems (NeurIPS), 2023.
10. **Keegan Harris**, Chara Podimata, and Zhiwei Steven Wu. *Strategic Apple Tasting*. Neural Information Processing Systems (NeurIPS), 2023.
11. **Keegan Harris***, Ioannis Anagnostides*, Gabriele Farina, Mikhail Khodak, Zhiwei Steven Wu, and Tuomas Sandholm. *Meta-Learning in Games*. International Conference on Learning Representations (ICLR), 2023.
12. **Keegan Harris**, Valerie Chen, Joon Sik Kim, Ameet Talwalkar, Hoda Heidari, and Zhiwei Steven Wu. *Bayesian Persuasion for Algorithmic Recourse*. Neural Information Processing Systems (NeurIPS), 2022.

*Denotes equal contribution and/or alphabetical ordering

13. **Keegan Harris**, Daniel Ngo*, Logan Stapleton*, Hoda Heidari, and Zhiwei Steven Wu. *Strategic Instrumental Variable Regression: Recovering Causal Relationships From Strategic Responses*. International Conference on Machine Learning (ICML), 2022.
14. **Keegan Harris**, Hoda Heidari, and Zhiwei Steven Wu. *Stateful Strategic Regression*. Neural Information Processing Systems (NeurIPS), 2021.

Preprints

1. **Keegan Harris**, Nicole Immorlica, Brendan Lucier, and Aleksandrs Slivkins. *Algorithmic Persuasion Through Simulation*. <https://arxiv.org/pdf/2311.18138>
2. **Keegan Harris**, Siddharth Prasad, and Asher Trockman. *In-Context Credit Assignment via the Core*. <https://arxiv.org/pdf/2605.06920>

Miscellaneous Writings

1. **Keegan Harris*** and Naveen Raman*. *Carnegie Mellon University at ICML 2025*. ML@CMU Blog, July 2025.
2. **Keegan Harris**. *Can ChatGPT Learn My Life From a Week of First-Person Video?* SIGBOVIK, 2025.
3. **Keegan Harris*** and Naveen Raman*. *Carnegie Mellon University at ICLR 2025*. ML@CMU Blog, April 2025.
4. **Keegan Harris*** and Naveen Raman*. *Carnegie Mellon University at NeurIPS 2024*. ML@CMU Blog, December 2024.
5. **Keegan Harris*** and Vasilis Syrgkanis*. *Causal Inference under Incentives: An Annotated Reading List*. SIGecom Exchanges, June 2024.
6. **Keegan Harris**. *Strategic Instrumental Variable Regression: Recovering Causal Relationships from Strategic Responses*. ML@CMU Blog, September 2021.

Part-Time Positions

Lux Capital	2025 – Present
<i>Graduate Associate</i>	
• Part of the investment team; focusing on AI, deep tech, and robotics startups • Responsibilities include conducting market research, participating in deal sourcing and diligence	
ML@CMU Blog	2024 – 2025
<i>Editor-in-Chief</i>	
• Under my leadership, the blog published 20 posts with a combined readership of 300,000+ views • Promoted to editor-in-chief after one month as editor; grew the editorial team from 2 to 7 editors	
Microsoft Research	Summer 2024
<i>Research Intern</i>	
• Conducted research with the Economics and Computation team, which resulted in a publication at UAI	
United States Innovative Technology	Summer 2024
<i>Technical Fellow</i>	
• Wrote USIT’s investment thesis on “AI in the Physical World”	
Microsoft Research	Summer 2023
<i>Research Intern</i>	
• Conducted research with the Economics and Computation team, which resulted in a publication at NeurIPS	
Vitable Health	Spring 2022
<i>Consultant</i>	
• Planned Vitable’s machine learning strategy and built their initial data collection pipeline	
United States Naval Research Laboratory	Summer 2019
<i>Intern</i>	
• Designed new planning algorithms for humanoid robot multitasking	
Penn State Applied Research Laboratory	2017 – 2018
<i>Distinguished Undergraduate Researcher</i>	
• Conducted applied machine learning research for DoD-funded projects	

Selected Talks

1. “*Should You Use Your Large Language Model to Explore or Exploit?*”
Carnegie Mellon University, AI Seminar (February 2025)
2. “*Regret Minimization in Stackelberg Games with Side Information*”
CUHK-Shenzhen, Global Young Scholars Forum (January 2025)
University of California San Diego, Rising Stars in Data Science Workshop (November 2024)
NeurIPS 2023, Workshop on Multi-Agent Security (December 2023, Spotlight Presentation)
3. “*The ML@CMU Blog: Perspectives After 5 Years*”
Carnegie Mellon University, CMU Libraries Forum for Student Publishing (September 2024)
4. “*Towards Bridging Causal Inference and Algorithmic Decision-Making*”
University of Illinois Urbana-Champaign, Machine Learning Seminar (September 2024)
Massachusetts Institute of Technology, Algorithms & Complexity Seminar (September 2024)
Microsoft Research Lab - New York City, Machine Learning Discussion Group (April 2024)
Capital One, Machine Learning Group (April 2024)
Columbia University, Machine Learning Group (January 2024)
5. “*Algorithmic Persuasion Through Simulation*”
Georgia Tech, ISyE Junior Researcher Workshop (April 2024)
Google Research NYC, Algorithms Seminar (January 2024)
Carnegie Mellon University, Economics Seminar (January 2024)
Drexel University, EconCS Seminar (October 2023)
6. “*Online Meta-Learning in Games and Bandits*”
Multi-Agent Learning Seminar (January 2024)
University of Miami, Industrial and Systems Engineering Seminar (September 2023)
7. “*Algorithmic Decision-Making using Panel Data*”
Federal Reserve Bank of Philadelphia, Conference on ML in Economics and Finance (December 2023)
University of Pennsylvania, Computer Science Theory Seminar (October 2023)
University of Miami, Computer Science Seminar (August 2023)
8. “*Algorithmic Decision-Making Under Incentives: Apple Tasting Feedback and Multiclass Learnability*”
Carnegie Mellon University, AI Seminar (September 2023)
Simons TOC For Fairness Seminar (May 2023)
9. “*Strategyproof Decision-Making in Panel Data Settings and Beyond*”
Online Causal Inference Seminar (May 2023)
10. “*Bayesian Persuasion for Algorithmic Recourse*”
3rd Symposium on the Foundations of Responsible Computing (June 2022)
AAMAS 2022, Learning with Strategic Agents Workshop (May 2022, Spotlight Presentation)
NeurIPS 2021, Workshop on Human and Machine Decisions (December 2021, Spotlight Presentation)
11. “*Strategic Instrumental Variable Regression: Recovering Causal Relationships From Strategic Responses*”
AAMAS 2022, Learning with Strategic Agents Workshop (May 2022, Spotlight Presentation)
ICML 2021, Workshop on Algorithmic Recourse (July 2021, Spotlight Presentation)
12. “*Stateful Strategic Regression*”
2nd Symposium on the Foundations of Responsible Computing (June 2021)

Academic Service

External Reviewer:

- ACM Conference on Economics & Computation (EC) 2022
- Conference on Neural Information Processing Systems (NeurIPS) 2021
- Innovations in Theoretical Computer Science (ITCS) 2024, 2022
- Symposium on Discrete Algorithms (SODA) 2022
- Symposium on Foundations of Responsible Computing (FORC) 2024
- Symposium on the Theory of Computation (STOC) 2024

Journal Reviewer:

- Games and Economic Behavior (GEB)
- Journal of Machine Learning Research (JMLR)

Professional Societies:

- INFORMS Annual Meeting Session Organizer: Online Marketplaces and Incentives 2024
- Organizer: [SIGecom Seminar Series](#) 2022

Program Committee:

- AAAI Conference on Artificial Intelligence (AAAI) 2025
- Conference on Artificial Intelligence and Statistics (AISTATS) 2025, 2024, 2023
- Conference on Learning Theory (COLT) 2025
- Conference on Neural Information Processing Systems (NeurIPS) 2026, 2025, 2024, 2023, 2022
- Conference on Uncertainty in Artificial Intelligence (UAI) 2026
- International Conference on Learning Representations (ICLR) 2026, 2025, 2024
- International Conference on Machine Learning (ICML) 2025, 2024, 2023, 2022
- The Web Conference (WWW) 2025

University Service:

- CMU Machine Learning Department Head Search Committee 2024
- CMU Machine Learning Department Speaking Skills Committee 2025, 2024
- CMU Machine Learning Ph.D. Admissions Committee 2024, 2023, 2022
- Editor: [ML@CMU Blog](#) 2024
- Editor-in-Chief: [ML@CMU Blog](#) 2025, 2024
- Mentor: CMU [Paths to AI Research](#) Mentoring Program 2024, 2023, 2022, 2021
- Mentor: CMU [SCS Graduate Application Support Program](#) 2024, 2023, 2022, 2020

Workshop Program Committee:

- Failure Modes of Agentic AI (ICML 2026)
- Formal Verification of Machine Learning (ICML 2023)
- Foundations of Reasoning in Language Models (NeurIPS 2025)
- In-Context Learning (ICML 2024)
- Learning and Decision Making with Strategic Feedback (NeurIPS 2021)
- Learning in Presence of Strategic Behavior (NeurIPS 2021)
- Mathematics of Modern Machine Learning (NeurIPS 2023)
- New Frontiers in Game-Theoretic Learning (ICML 2026)
- Privacy-Preserving Artificial Intelligence (AAAI 2024)
- Trustworthy and Socially Responsible Machine Learning (NeurIPS 2022)

Teaching

Guest Lecturer:

- Foundations of Learning, Game Theory, and Their Connections (Spring 2025). Instructor: [Nina Balcan](#)
- Foundations of Operations Management (Fall 2023). Instructor: [Sridhar Tayur](#)
- Selected Topics in Applied Statistics (Spring 2024). Instructors: [Raaz Dwivedi](#) and [Kyra Gan](#)

Teaching Assistant:

- Advanced Topics in Machine Learning Theory (Fall 2022). Instructor: [Nina Balcan](#)
- Operations Management (Fall 2023). Instructor: [Sridhar Tayur](#)
- Service Management (Fall 2023). Instructor: [Sridhar Tayur](#)